

Nithin Raghavan

(678) 200-5839 — rnithin@berkeley.edu — rnithin1 (Github) — linkedin.com/in/nithinraghavan

EDUCATION

University of California, Berkeley (Class of 2020)

Aug 2017 – present

Computer Science *Bachelor of Arts*, Applied Mathematics *Bachelor of Arts*

(GPA: 3.67)

- CS61B: Data Structures
- CS170: Efficient Algorithms
- EE127: Optimization Models and Applications
- Math 128a: Numerical Analysis
- Math 126: Partial Differential Equations
- CS189: Introduction to Machine Learning

Georgia Institute of Technology

Aug 2015 – May 2017

Courses Taken in High School

- Math 3012: Applied Combinatorics
- Math 2803: Number Theory and Cryptography

EXPERIENCE

→ Samsung Advanced Computing Lab

May 2019 – Present

- Conducted extensive research on deep learning usecases and models as part of Samsung's GPU team
- Conducted extensive research on the potential routes of optimization and quantization of deep learning models such as MobileNet, R-FCN, SRCNN and ESRGAN as part of Samsung's GPU team
- Wrote and implemented OpenCL and OpenGL code
- Researched the graphics pipeline and became acquainted with AMD's compute and graphics architecture, and wrote + ran 2D register-blocked GEMM kernels with increased WPT in OpenCL on AMD architecture
- Wrote and trained two neural networks; the first to perform ambient occlusion on complex OpenGL-rendered scenes, and the second to convert a flat-rendered scene to a lifelike, physically based rendered one

→ Mobile Sensing Lab, UC Berkeley

Oct 2018 – Present

- Currently writing code implementing a parallelized Frank-Wolfe algorithm for dynamic traffic assignment in C++/CUDA using contraction hierarchies
- Helping research the impact of different optimization models of routing behaviour on the Waze problem

→ RISE Lab, UC Berkeley

Jun 2018 – Dec 2018

- Designed and implemented a data visualization tool for Jupyter Notebook for hyperparameter optimization for Cirrus, a serverless machine learning framework
- Helped write code involving AWS Lambdas for model primitives such as logistic regression

→ IBM Almaden Research Center, Machine Learning Laboratory

Jul 2017 – Aug 2017

- Trained an artificial neural network with visual question answering abilities on Stanford's CLEVR dataset with 70% overall accuracy
- Implemented sequence autoencoders, CNNs and LSTMs with Tensorflow and Keras

PROJECTS

→ Resource-Provisioning GPU Server

Dec 2017 – present

- Developed a Python-based shell to automate on-demand request processing and resource provisioning in a GPU + CPU cluster
- Collaborated on a team to create a program that utilizes Slurm for cluster management and deploys tasks in Docker containers

→ Software Renderer

Jul 2019

- Developed a software-based rasterizer and renderer with pixel and vertex shader support in C++
- Capable of barycentric interpolation, backface culling and block-based rasterization

→ LASSO/Wavelet Based Compressed Sensing

Jul 2019

- Computes LASSO on the matrix-vector product representation of the discrete wavelet transform of an input signal with orthogonal Daubechies wavelets
- Can lossily compress audio/images to any amount or preprocess them for ML training purposes

→ TaxiFindMe

Apr 2018

- Routing web app that helps New Yorkers find the best spot to minimize taxi waiting time, taking into account travel time and time of day

- Preprocessed 20 million entry taxi dataset with k-means machine learning algorithm; for querying, KNN is run from an input location to find nearest cluster. Frontend employs Django
- Reduced query time up to 94% from the naive implementation

→ **ShirtMapper**

Jan 2018

- App that resizes images of custom shirts and maps them onto people
- Utilizes OpenCV and Scipy, and uses Haar classifiers for edge detection; frontend employs React Native

SKILLS

Awards: Exploravision National Contest

2016

- Wrote a paper proposing blockchain's potential link to autonomous vehicles, and won honorable mention.

Models/Algorithms: Regression/classification (ridge, logistic, SVM, decision trees, OLS), PCA/SVD, ensemble learning, k-means, deep learning (CNNs, LSTMs, GANs), Frank-Wolfe

Frameworks/Softwares: Numpy, Scipy, Pytorch, OpenCV, Docker, Slurm, d3js, OpenCL, OpenGL

Programming Languages: Python, Java, C, C++, C#, Bash, Latex, SQL, JavaScript, Matlab

Operating Systems: Unix-like systems (Linux, FreeBSD, Mac OS X), Windows

Certifications: Android Development (University of Maryland through Coursera)